

Part 3

Inferences about μ
Confidence Intervals
Type I and II Error
p-value
t-distribution

Inferences about μ

- One of the major objectives of statistics is to make inferences about the distribution of the elements in a *population* based on information contained in a *sample*.
- Numerical summaries that characterize the population distribution are called *parameters*.
- The *population mean* μ and *population variance* σ^2 are two important parameters. Others are median, range, mode, etc.

Inferences about μ - cont.

- Methods for **making inferences** are basically designed to answer one of two types of questions:
 - Approximately what is the value of the parameter?
or
 - Is the value of the parameter less than (say) 6?
- Statisticians answer the first question by **estimating** the parameter using the sample.
- The second case might require a **test of a hypothesis**.

Point and Interval Estimation of μ when σ is known and n is large

- Point estimation of μ does not require that σ be known or a large n
- The point estimate of μ is the sample mean $\bar{y}$
- Point estimates by themselves do not tell how much $\bar{y}$ might differ from μ , that is, the accuracy or precision of the estimate
- A measure of accuracy is the difference between sample mean $\bar{y}$ and population mean μ is called **sampling error**

A Confidence Interval for μ

- An interval estimate, called a *confidence interval*, incorporates information about the amount of sampling error in $\bar{y}$
- A *confidence interval* for μ takes the form

$$(\bar{y} - E, \bar{y} + E),$$

for a number E

- An associated number called the *confidence coefficient* helps assess how likely it is for μ to be in the interval.

A Confidence Interval for μ - cont.

- To derive the specific form of the confidence interval for μ , for the case **when σ is known and n is large**, the CLT result must be used.
- By the CLT, $Z = (\bar{Y} - \mu)/\sigma_{\bar{Y}}$ is approximately $\sim N(0, 1)$.
- Let Z have a $N(0, 1)$ distribution (exactly).
- Let $z_{\alpha/2}$ denote the $1 - \alpha/2$ quantile of the standard normal distribution, for a given number α , $0 < \alpha < 1$.
- Then the following probability statement is true:

$$P\left(-z_{\alpha/2} \leq \frac{(\bar{Y} - \mu)}{\sigma_{\bar{Y}}} \leq z_{\alpha/2}\right) = 1 - \alpha$$

A Confidence Interval for μ - cont.

- Manipulating the inequalities, without changing values, we have

$$P(\sigma_{\bar{Y}} z_{\alpha/2} \geq \mu - \bar{Y} \geq -\sigma_{\bar{Y}} z_{\alpha/2}) = 1 - \alpha$$

$$P(\bar{Y} - \sigma_{\bar{Y}} z_{\alpha/2} \leq \mu \leq \bar{Y} + \sigma_{\bar{Y}} z_{\alpha/2}) = 1 - \alpha$$

- If, for example, $\alpha = 0.05$ then

$$P(\bar{Y} - \frac{\sigma}{\sqrt{n}} z_{0.025} \leq \mu \leq \bar{Y} + \frac{\sigma}{\sqrt{n}} z_{0.025}) = 0.95$$

and, since $z_{0.025} = 1.96$,

$$P(\bar{Y} - 1.96 \frac{\sigma}{\sqrt{n}} \leq \mu \leq \bar{Y} + 1.96 \frac{\sigma}{\sqrt{n}}) = 0.95$$

A Confidence Interval for μ - cont.

- Using this statement, a $(1-\alpha)100\%$ confidence interval for μ , with confidence coefficient $1 - \alpha$ can be calculated using a random sample of data $y_1, y_2, \dots, y_n$
- It is usually written in the form

$$\left(\bar{y} - \frac{\sigma}{\sqrt{n}} z_{\alpha/2}, \bar{y} + \frac{\sigma}{\sqrt{n}} z_{\alpha/2} \right) .$$

- **Example:** Suppose $n = 36$, $\sigma = 12$, $\bar{y} = 24.8$

And for $\alpha = 0.05$, $z_{\alpha/2} \equiv z_{0.025} = 1.96$

A Confidence Interval for μ - cont.

- Thus a 95% C.I. for μ is

$$\left(24.8 - \frac{12}{\sqrt{36}} \times 1.96, \quad 24.8 + \frac{12}{\sqrt{36}} \times 1.96 \right)$$

- The interval for μ is: (20.88, 28.72) with confidence coefficient 0.95.
- We might say that we are *95% confident* that the population mean μ is between 20.88 and 28.72.
- But what do we actually mean when we say that we are *95% confident*?

Interpretation of a Confidence Interval

- Before the sample is drawn, the probability is $(1 - \alpha)$ that the **random interval** $(\bar{Y} - \frac{\sigma}{\sqrt{n}} z_{\alpha/2}, \bar{Y} + \frac{\sigma}{\sqrt{n}} z_{\alpha/2})$ will contain μ (because $\bar{Y}$ is a random variable).
- However, once the sample is drawn, and $\bar{y}$ is calculated, the interval ceases to be random. It is a numerical interval $(\bar{y} - \frac{\sigma}{\sqrt{n}} z_{\alpha/2}, \bar{y} + \frac{\sigma}{\sqrt{n}} z_{\alpha/2})$, calculated specifically for the drawn sample; thus we cannot associate a **probability** with it.
- We may say that the process which led us to this interval will, on the average, produce an interval containing μ , $100(1 - \alpha)\%$ of the time.

Interpretation of a Confidence Interval - cont.

- *It is not correct to say that the interval $(\bar{y} - \frac{\sigma}{\sqrt{n}} z_{\alpha/2}, \bar{y} + \frac{\sigma}{\sqrt{n}} z_{\alpha/2})$ contains μ with a specified probability.*
- If the process of sampling is repeated many times and 100(1- α)% confidence intervals calculated for each sample, then we are confident that 100(1- α)% of those intervals will contain μ .

Notes:

- As the confidence coefficient increases the confidence interval becomes wider. A wider interval estimates μ less precisely. Thus a 99% confidence interval is **less accurate** than a 95% confidence interval but one has **more confidence** that μ is in the first interval.
- Also note that increasing the sample size n results in a narrower interval (more accurate estimate) for the same confidence coefficient.

Example:

The caffeine content (in mg) was examined for a random sample of 50 cups of black coffee dispensed by a new machine. The mean and standard deviation were 110 mg and 7.1 mg, respectively. Use these data to construct a 98% confidence interval for μ , the mean caffeine content for cups dispensed by the machine.

As the sample size is large we can use the CLT and also approximate σ , the sample standard deviation of the population by $s = 7.1$

$$n = 50 \quad \bar{y} = 110, \quad \sigma = 7.1, \quad \alpha = .02, \quad z_{0.01} = 2.33$$

Example - cont.

$$\left[\bar{y} - z_{0.01} \frac{\sigma}{\sqrt{n}}, \bar{y} + z_{0.01} \frac{\sigma}{\sqrt{n}} \right]$$
$$\left[110 - (2.33) \left(\frac{7.1}{\sqrt{50}} \right), 110 + (2.33) \left(\frac{7.1}{\sqrt{50}} \right) \right]$$
$$(107.66, 112.34)$$

We are 98% confident that the mean caffeine content for cups dispensed by the machine is between 107.66 and 112.34 mg.

Choosing the Sample Size

- The **width** of a confidence interval is $(2 \times \sigma \times z_{\alpha/2})/\sqrt{n}$.
- The width of the confidence interval can be made smaller by changing to a larger α or increasing sample size n .
- Let us consider selection of n to achieve a desired width $2 \cdot E$ for a fixed α
- We want E to at least equal to $\sigma \times z_{\alpha/2}/\sqrt{n}$
- Thus we need to select a sample so that

$$n \geq \frac{(z_{\alpha/2})^2 \sigma^2}{E^2}$$

Choosing the Sample Size - cont.

- This will result in an interval no wider than $(\bar{y} - E, \bar{y} + E)$.

Example:

Given $\sigma = 12$, $\alpha = 0.05$, $z_{0.025} = 1.96$

What sample size n will give an interval no wider than 5.6?

We set $E = 2.8$ so

$$n \geq \frac{(1.96)^2 \cdot 144}{(2.8)^2} = 70.56$$

Thus the experimenter must choose a sample size $n = 71$ at least.

Choosing the Sample Size - cont.

Would the sample size $n = 71$ ensure the width to be ≤ 5.6 if the population has a larger variance? Would $n = 71$ be enough?

The answer is **NO** since the formula

$$\frac{(z_{\alpha/2})^2 \sigma^2}{E^2}$$

involves the population variance σ^2 .

Precision in estimation depends on both α and σ^2 . If the variance of population elements is very small, i.e., the elements are tightly clustered about the population mean μ , then only a small sample is needed for the estimate $\bar{y}$ to be very near μ .

Statistical Tests for μ

- Estimation (either point or interval estimation) was used to help answer a question like “Approximately what is the value of μ ?”
- The other kind of question mentioned earlier is “Is it likely that μ is less than (or greater than) the value μ_0 (a pre-determined value)?”.
- A **Test of Hypothesis** is used to answer this kind of question. As in estimation, the sample mean $\bar{y}$ of a random sample of n elements from the population is used to answer this question.

Statistical Tests for μ - cont.

- Every test of hypothesis features
 - (a) a **Null Hypothesis** H_0 which describes a characteristic of the population (it is believed to be) as it currently exists,
 - (b) a **Research Hypothesis** (or **Alternative Hypothesis**) H_a , which is a proposal about this characteristic, by the person(s) conducting the statistical study.
- The idea is that the null hypothesis is presumed to hold unless there is overwhelming evidence in the data in support of the research hypothesis.

Example:

- To determine whether the mean yield per acre (in bushels), μ , of a variety of soybeans increased the current year over the last two years when μ is believed to be 520 bushels per acre, the following might be tested.


$$H_0 : \mu \leq 520 \quad \text{vs.} \quad H_a : \mu > 520$$

- The fact that μ may equal a specific value is always included in the null hypothesis. The decision to state whether the data supports the research hypothesis or not is based on a quantity computed from the data called the **test statistic**.

- If $\bar{y}$ is in the rejection region then reject H_0 and say the evidence favors H_a .
- Basically, the test amounts to computing $\bar{y}$ and looking at its value relative to μ_0 and the μ values in H_a .
- For example, in the above example if $\bar{y} < 520$ we will say **there is not sufficient evidence to reject H_0** .
- Even if $\bar{y} > 520$, we might still say **there is not sufficient evidence to reject H_0**
- This is ok **as long as $\bar{y}$ is not too much greater than 520**.
- **"How much is too large? "** To decide this, use the probability distribution of $\bar{Y}$. Begin by picking a small probability α like $\alpha = 0.05$, or $\alpha = 0.01$, or $\alpha = 0.001$.

- Then reject H_0 only when the probability of obtaining a value of $\bar{y}$ larger than the observed value $\bar{y}$ is $\leq \alpha$ when μ is indeed ≤ 520 , i.e., when H_0 is true.
- That is, if H_0 is true, the chance of observing a sample that results in a $\bar{y}$ as large as the one calculated should be very small.

FIGURE 5.7
that H_0 is true,
values of $\bar{y}$ are
the upper tail.



How to determine the rejection region

- By the CLT, $\bar{Y}$ is approximately a normal random variable.
- We will ignore the approximation and assume $\bar{Y} \sim N(\mu, \frac{\sigma^2}{n})$.
- Then, for a given α ,

$$P\left(\frac{\bar{Y} - \mu}{\sigma/\sqrt{n}} > z_\alpha\right) = \alpha$$

where z_α is the $1 - \alpha$ quantile of the standard normal distribution.

- So when $H_0 : \mu \leq \mu_0$ is true ,

$$P\left(\frac{\bar{Y} - \mu_0}{\sigma/\sqrt{n}} > z_\alpha\right) \leq \alpha$$

How to determine the rejection region - cont.

- That is, $P(\bar{Y} > \mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}) \leq \alpha$. This suggests that we reject $H_0 : \mu \leq \mu_0$ in favor of, say $H_a : \mu > \mu_0$ when

$$\bar{y} > \mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$$

(i.e., $\bar{y}$ is **too much** when $\bar{y}$ exceeds the number $\mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$)

- All possible $\bar{y}$ values satisfying $\bar{y} > \mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$ constitute the **rejection region**.
- Instead of comparing $\bar{y}$ to $\mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$, it is easier to first calculate

$$z_c = \frac{\bar{y} - \mu_0}{\sigma / \sqrt{n}}$$

How to determine the rejection region - cont.

- We see that comparing $\bar{y}$ to $\mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$ is the same as comparing z_c to z_α .
- That is, we reject the null hypothesis if $z_c > z_\alpha$ which is the same thing as doing so if $\bar{y} > \mu_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$
- This quantity z is called the **test statistic** and its value can be calculated using the data and the value μ_0 specified in the null hypothesis H_0 .
- The notation z_c is used to denote the “computed value” of the test statistic, that is when we plug in the observed value of $\bar{y}$ and obtain a numerical value for z .

Type I Error

- The α corresponds to the probability that the null hypothesis is rejected when actually it is true and is called the probability of committing a Type I error.
- Since the experimenter selects the value of α used in the test procedure, she is able to specify or control the Type I error rate or “how much” of this type of error is permitted in the testing procedure.

Example

Suppose from a sample of 36 1-acre plots, the yield of corn this year was measured and $\bar{y} = 573$ and $s = 124$ calculated. Can we conclude that the mean yield of corn for all farms exceeded 520 bushels/acre this year ? Here we are going to assume that σ can be approximated by s . Use $\alpha = .025$

Solution to the Example

Set-up five parts of the testing procedure:

- 1 $H_0 : \mu \leq 520$
- 2 $H_a : \mu > 520$
- 3 T.S.: $z = \frac{\bar{y} - \mu_0}{\sigma / \sqrt{n}}$
- 4 R.R: Reject H_0 : in favor of H_a : if $z > z_{.025}$
i.e. the R.R. is $z > 1.96$ since $z_{.025} = 1.96$
- 5 Compute the observed value of the test statistic:
 $z_c = \frac{573 - 520}{124 / \sqrt{36}} = 2.56$

Decision:

Since $z_c > 1.96$, z_c is in the R.R. Thus $H_0 : \mu \leq 520$ is rejected in favor of the research hypothesis $H_a : \mu > 520$. It is concluded that the mean yield this year exceeds 520 bushels/acre.

Note that the above test procedure is equivalent to determining that the observed value for $\bar{y}$ lies more than 1.96 standard deviations above the mean $\mu_0 = 520$.

Summary of Test Procedures - Assume n large, and σ known

Hypotheses:

Case 1: $H_0 : \mu \leq \mu_0$ vs. $H_a : \mu > \mu_0$ (*right-tailed test*)

Case 2: $H_0 : \mu \geq \mu_0$ vs. $H_a : \mu < \mu_0$ (*left-tailed test*)

Case 3: $H_0 : \mu = \mu_0$ vs. $H_a : \mu \neq \mu_0$ (*two-tailed test*)

T.S:
$$z_c = \frac{\bar{y} - \mu_0}{\sigma / \sqrt{n}}$$

R.R: For Type I error probability of α :

Case 1: Reject H_0 : if $z \geq z_\alpha$

Case 2: Reject H_0 : if $z \leq -z_\alpha$

Case 3: Reject H_0 : if $|z| \geq z_{\alpha/2}$

Decision:


If the computed value of the test statistic, z_c , is in the R.R., then we will reject the null hypothesis H_0 : at the specified α value. Otherwise, we say we fail to reject H_0 : at the specified α value.

Example A corporation maintains a large fleet of company cars for its salespeople. To check the average number of miles driven per month per car, a random sample of $n = 40$ cars is examined. The mean and standard deviation for the sample are 2,752 miles and 350 miles, respectively. Records for previous years indicate that the average number of miles driven per car per month was 2,600. Use the sample data to test the research hypothesis that the current mean μ differs from 2,600. Set $\alpha = .05$ and assume that σ can be replaced by s .

Solution

The null hypothesis for this statistical test is $H_0 : \mu = 2,600$ and the research hypothesis is $H_a : \mu \neq 2,600$. Using $\alpha = .05$, the two-tailed rejection region for this test is $|z| > z_{.025}$ or $|z| > 1.96$ and is located as shown below.

FIGURE 5.11
Region for H_a
when $\alpha = .05$



Solution - cont.

To determine how many standard errors our test statistics $\bar{y}$ lies away from $\mu = 2,600$, compute

$$z_c = \frac{\bar{y} - \mu_0}{\sigma/\sqrt{n}} = \frac{2,752 - 2,600}{350/\sqrt{40}} = 2.75.$$

Thus $|z_c| = 2.75 > 1.96$ and therefore we reject $H_0 : \mu = 2,600$ at $\alpha = .05$. It follows that the observed value for $\bar{y}$ lies more than 1.96 standard errors above the mean $\mu = 2,600$, so we reject the null hypothesis in favor of the alternative $H_a : \mu \neq 2,600$.

Since $\bar{y} > 2,600$ we conclude that the mean number of miles driven is greater than 2,600.

Level of Significance or the p-value of a Statistical Test

- As an alternative to the formal test where one uses a rejection region based on a specified the Type I error rate α , many researchers compute and report the **level of significance** or the **p-value** for the test.
- This is the probability, when H_0 is true, of observing a statistic as extreme as the one actually observed. Here “extreme” is means “large” or “small” according to the alternative hypothesis H_a .

Level of Significance or the p-value of a Statistical Test - cont.

- Specifically, for a given σ^2 and μ_0 we find the p-value by:
 - First computing the test statistic, z_c ,

$$z_c = (\bar{y} - \mu_0)/(\sigma/\sqrt{n})$$

- If $H_a : \mu > \mu_0$, then $p = P(Z > z_c)$.
 - If $H_a : \mu < \mu_0$, then $p = P(Z < z_c)$.
 - If $H_a : \mu \neq \mu_0$, then $p = 2P(Z > |z_c|)$.
- A p-value smaller than the pre-specified α value is evidence in favor of rejecting H_0 .

Example:

- In the earlier example we tested $H_0 : \mu \leq 380$ vs. $H_a : \mu > 380$
Calculating the z -statistic,

$$z_c = \frac{\bar{y} - 380}{\sigma/\sqrt{n}} = \frac{390 - 380}{35.2/\sqrt{50}} = 2.01$$

- The **level of significance**, or **p-value** for this test is

$$p = P(Z > 2.01) = 1 - P(Z < 2.01) = 0.0222$$

- We **fail to reject** $H_0 : \mu \leq 380$ at $\alpha = .01$ since p-value is not less than .01

Inferences About μ when σ is unknown

- For large sample sizes, it follows from the CLT that $\bar{Y}$ is approximately Normally distributed i.e., $\bar{Y} \sim N(\mu, \frac{\sigma^2}{n})$.
- It follows that the random variable

$$T_{n-1} = (\bar{Y} - \mu) / (S / \sqrt{n})$$

has approximately the **Student's t distribution** with $n - 1$ **degrees of freedom**.

- Here $Y_1, Y_2, \dots, Y_n$ are sampling random variables and $\bar{Y} = \frac{1}{n} \sum_i Y_i$ is thus a random variable.

Inferences About μ when σ is unknown - cont.

- The denominator of T_{n-1} is the random variable

$$S = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2}.$$

- If sampling from a Normal distribution, the sample mean $\bar{Y}$ has a Normal distribution (exactly) and therefore $(\bar{Y} - \mu)/(S/\sqrt{n})$ will have a Student's t distribution (exactly), regardless of sample size.


- Until now, we have assumed that n is large and σ is known, or n is large and n is large enough also to use s as an approximation to σ .
- Since we used the CLT to derive confidence limits and tests of hypotheses, the exact nature of the sampled population was not required to be specified.
- Now we require that the population distribution to be Normal. We don't need to know σ or have a large sample, i.e. n to be large.
- In this situation, for any n , $(\bar{Y} - \mu)/(S/\sqrt{n})$ will have a Student's t-distribution with d.f. = $n - 1$.
- Using this fact we can obtain confidence intervals and conduct tests of hypotheses even though σ is unknown.
- Note the difference between Z and T_{n-1} is that the parameter σ is in the denominator of Z .
- That is, the point estimator of σ , S is in the denominator of T_{n-1} .

Properties of the Student's t distribution

- There are many t -distributions each specified by a single parameter called **degrees of freedom (df)**.
- Like the standard normal population, the distribution is **symmetric** about 0 and has mean equal to 0.
- The t -distribution has variance $df/(df - 1)$, and hence is more variable than the standard normal distribution which has variance equal to 1.
- We say that the t -distribution has **heavier tails** than the standard normal distribution.
- As the degrees of freedom df increases, the t -distribution approaches that of the standard normal distribution.
- Thus as the sample size n increases the distribution of the T_{n-1} random variable approaches the standard normal distribution.

Properties of the Student's t distribution cont.

FIGURE 5.16
The t distribution
approaches a standard
normal distribution



The t table gives quantiles 0.90, 0.95, 0.975, 0.99, 0.995, and 0.999 for t distributions with selected df .

Examples: For $df = 10$

$$P(T_{10} > 1.372) = 0.10$$

$$P(T_{10} < 1.372) = 0.90$$

$$P(T_{10} > 1.812) = 0.05$$

If $\alpha = 0.05$, then t_α satisfies $P(T_{10} > t_{.05}) = 0.05$ and from the t table, $t_{.05} = 1.812$. If $\alpha = 0.05$, then $t_{\alpha/2} \equiv t_{0.025} = 2.228$

Comparison of Normal and t Quantiles

t_α with indicated df

α	z_α	10	20	30	40	60	240
0.1	1.28	1.37	1.32	1.31	1.30	1.30	1.28
0.05	1.65	1.81	1.72	1.70	1.68	1.67	1.65
0.01	2.33	2.76	2.53	2.46	2.42	2.39	2.34

Confidence Interval for μ based on the t-distribution

- Let the sample size be n , and for α be a specified value, say, for e.g. .05.
- A $(1 - \alpha)100\%$ confidence interval for μ is given by

$$\left(\bar{y} - t_{\alpha/2} \frac{s}{\sqrt{n}}, \quad \bar{y} + t_{\alpha/2} \frac{s}{\sqrt{n}} \right)$$

- This may also be written as $(\bar{y} - t_{\alpha/2}s_{\bar{y}}, \quad \bar{y} + t_{\alpha/2}s_{\bar{y}})$
- Here, $t_{\alpha/2}$ is the $1 - \alpha/2$ percentile of the t-distribution with $n - 1$ degrees of freedom and $s_{\bar{y}}$ is the **standard error of the mean**.

EXAMPLE 5.17

An airline wants to evaluate the depth perception of its pilots over the age of fifty. A random sample of $n = 14$ airline pilots over the age of fifty are asked to judge the distance between two markers placed 20 feet apart at the opposite end of the laboratory. The sample data listed here are the pilots' error (recorded in feet) in judging the distance.


2.7 2.4 1.9 2.6 2.4 1.9 2.3
2.2 2.5 2.3 1.8 2.5 2.0 2.2

Use the sample data to place a 95% confidence interval on μ , the average error in depth perception for the company's pilots over the age of fifty.

SOLUTION Before setting up a 95% confidence interval on μ , we must first assess the normality assumption by plotting the data in a normal probability plot or a boxplot. Figure 5.20 is a boxplot of the 14 data values. The median line is near the center of the box, the right and left whiskers are approximately the same length, and there are no outliers. The data appear to be a sample from a normal distribution. Thus, it is appropriate to construct the confidence interval based on the t distribution. You can verify that

$$\bar{y} = 2.26 \text{ and } s = .28$$

FIGURE 5.20
Boxplot of distance (with
95% confidence interval for
the mean)



Referring to Table 2 in the Appendix, the t -value corresponding to $\alpha = .025$ and $df = 13$ is 2.160. Hence, the 95% confidence interval for μ is

$$\bar{y} \pm t_{\alpha/2} \frac{s}{\sqrt{n}} \text{ or } 2.26 \pm 2.160 \frac{.28}{\sqrt{14}}$$

which is the interval $2.26 \pm .16$, or 2.10 to 2.42. Thus, we are 95% confident that the average error in the pilots' judgment of the distance is between 2.10 and 2.42 feet.

Test Procedures based on the t-distribution:

Hypotheses:

Case 1: $H_0 : \mu \leq \mu_0$ vs. $H_a : \mu > \mu_0$ (right-tailed test)

Case 2: $H_0 : \mu \geq \mu_0$ vs. $H_a : \mu < \mu_0$ (left-tailed test)

Case 3: $H_0 : \mu = \mu_0$ vs. $H_a : \mu \neq \mu_0$ (two-tailed test)

T.S: $t_c = \frac{\bar{y} - \mu_0}{s/\sqrt{n}}$

R.R: For Type I error probability of α :

Case 1: Reject H_0 : if $t \geq t_{\alpha, (n-1)}$

Case 2: Reject H_0 : if $t \leq -t_{\alpha, (n-1)}$

Case 3: Reject H_0 : if $|t| \geq t_{\alpha/2, (n-1)}$

Test Procedures based on the t-distribution - cont.:

Level of Significance (p-value):

Case 1: $p = P(T_{n-1} > t_c)$.

Case 2: $p = P(T_{n-1} < t_c)$.

Case 3: $p = 2P(T_{n-1} > |t_c|)$.

Exercise

A massive multistate outbreak of food-borne illness was attributed to *Salmonella enteritidis*. Epidemiologists determined that the source of the illness was ice cream. They sampled nine production runs from the company that produced the ice cream to determine the level of *Salmonella enteritidis* in the ice cream.

Exercise - cont.

These levels (MPN/g) are as follows

.593 .142 .329 .691 .231 .793 .519 .392 .418

Use the data to determine whether the mean level of *Salmonella enteritidis* in the ice cream is greater than .3 MPN/g with $\alpha = .01$


Solution:

Need to test $H_0 : \mu \leq .3$ vs. $H_a : \mu > .3$

Because of the small sample size, we need to examine whether the data have been sampled from a normal distribution. To do this a **normal probability plot** is a good tool.

Exercise - cont.

FIGURE 5.18
Normal probability plot for
Salmonella data



From the data $\bar{y} = .456$ and $s = .2128$ are computed, giving

$$t_c = \frac{\bar{y} - \mu_0}{s/\sqrt{n}} = \frac{.456 - .3}{.2128/\sqrt{9}} = 2.21$$

Because for the one-tailed test we need $t_{\alpha, (n-1)}$;
we look up $t_{.01}$ with $df = 9 - 1 = 8$. It is 2.896

Thus the rejection region is: $t > 2.896$

Exercise - cont.

Since $t_c = 2.214$ does not exceed 2.896 it is not in the R.R.

Thus, there is insufficient evidence in the data to reject H_0 i.e to say that the mean level of *Salmonella enteritidis* exceeds the dangerous level of .3 MPN/g.

The **p-value** to be computed is $P(T_8 > 2.21)$

To calculate this exactly using the t-table is not possible since it is tabulated for only a few values of a .

However, we can bound the p-value by noting that, for ($df = 8$), 2.21 lies between 1.86 and 2.306.

This gives $.025 < \text{p-value} < .05$ showing that the **p-value is not less than our α of .01**. Thus we fail to reject H_0 .

OC Curve and the Power of a test

The probabilities of the four possible outcomes of a statistical test are:

Decision	Null Hypothesis	
	True	False
Reject H_0	Type I Error α	Correct Decision $1 - \beta$
Fail to reject H_0	Correct Decision $1 - \alpha$	Type II Error β

- The probabilities of Type I and II error are α and β , respectively.

Type I Error and Type II Error

- The experimenter can control only the Type I error probability. We do this by specifying an α for an experiment (before the data values are measured).
- When we are testing an hypothesis about μ using α for the test, the value of Type II error probability β depends on the actual value μ which is not known.
- This is because β is the probability of incorrectly failing to reject H_0 when H_a is true.
- When H_a is true, the actual value of μ may be any value under H_a , i.e, a value not specified under H_0 . Since this value of μ is unknown, β cannot be calculated.


Type I Error and Type II Error

- The implication of this is that when, based on a test, we find that we cannot reject H_0 , *we will not say that we accept H_0 .*
- Because if we say we accept H_0 and β turns out to be large, then the probability is large that we will be committing a Type II error.
- Thus the correct way to state the decision is to say, that **we fail to reject H_0 .**
- In practice, β probabilities for several choices of μ (call them $\beta(\mu)$) are calculated and plotted in a graph called the **OC curve.**
- The OC curve can be used to read-off the **Type II error probability** for a specified set of values sample size n and α .

OC Curve

FIGURE 5.14

(a) OC curve for $\alpha = .05$, $n = 10, 18, 25$. (b) OC curve for $n = 25$, $\alpha = .05, .01, .001$



The figure shows the OC curve for the test of $H_0 : \mu \leq 84$, $H_a : \mu > 84$ for a population with $\sigma = 1.4$. For example, for $\alpha = .05$, $n = 10$, it can be seen that $\beta(84.8) \approx .4$ and that $\beta(84.8)$ decreases as sample size goes from 10 to 25. A conclusion that can be made about this test from the OC curve is that Type II error probability would be $< .1$ for an actual $\mu > 84.8$ for $n = 25$.

Power Curve

- Another quantity that may be calculated for a test procedure for a specified value of μ is called the **Power of the test** and is defined as $1 - \beta(\mu)$.
- The corresponding plot of power against a set of μ values is called the **power curve**.


Power of the Test

- By definition, power of a test is the probability of rejecting H_0 for a specified value of μ under H_a .
- In practice, tests are designed to have large power for some μ values of interest so that they have small Type II error probabilities.
- We can relate to this idea by thinking of a test as having very good power if it has a very good chance of detecting whether a change in μ has actually occurred.
- This is usually done by selecting a sample size n to be used for the experiment so that the desired power is achieved for a specified μ and α .

Using Type II Error Probability β Curves

- Consider the Salmonella example again. We have $n = 9$, and $\alpha = .01$. Thus $df = 8$ and we estimate $\sigma \approx .25$
- We can compute the values of d for several values of μ_a . Then we'll read β for those values of d using the graph on next slide for the curves for $\alpha = .01$.
- As an example, for $\mu_a = .45$, $d = \frac{|\mu_a - \mu_0|}{\sigma} = \frac{|.45 - .3|}{.25} = .6$
- Corresponding to $d = .6$ on the horizontal axis, using the curve for $df = 8$ we see that $\beta(.45) = .79$, approx.
- Similarly, for $\mu_a = .45$ $d = 1.0$, and thus $\beta(.55) = .43$, approx. We can construct a table as shown (next slide).

Using Type II Error Probability β Curves



μ_a	.35	.4	.45	.5	.55	.6	.65	.7	.75
d	.2	.4	.6	.8	1.0	1.2	1.4	1.6	1.8
$\beta(\mu_a)$	.97	.91	.79	.63	.43	.26	.13	.05	.02

Departures from Normality

- When $\bar{Y}$ is a Normal random variable, i.e., when sampling from a Normal population with mean μ_0 , $\frac{\bar{Y} - \mu_0}{S/\sqrt{n}}$ is a T_{n-1} random variable.
- This is a theoretical fact. In practice, however, we never sample exactly from a Normal population, so $(\bar{Y} - \mu_0)/S/\sqrt{n}$ will be only approximately T_{n-1} .
- How much effect can this have on C.I.'s and tests we construct?
- For symmetrically distributed populations and not too small n there is little to worry about.
- For highly skewed population distributions the approximation can be terrible especially for small n .

TABLE 5.6
Level and power values
for t test

Population Distribution	$n = 10$				$n = 15$				$n = 20$			
	Shift d				Shift d				Shift d			
	0	.2	.6	.8	0	.2	.6	.8	0	.2	.6	.8
Normal	.05	.145	.543	.754	.05	.182	.714	.903	.05	.217	.827	.964
Heavy tailed	.035	.104	.371	.510	.049	.115	.456	.648	.045	.163	.554	.736
Light skewness	.025	.079	.437	.672	.037	.129	.614	.864	.041	.159	.762	.935
Heavy skewness	.007	.055	.277	.463	.006	.078	.515	.733	.011	.104	.658	.873

- It is recommended that one look at a boxplot, Normal plot, and/or other graphics to see whether severe skewness of the sampling population is indicated. If not, proceed to use the t -distribution. If yes, one can use a nonparametric procedure or use a transformation. We will look at some of these later.

Using Confidence Intervals to Test Hypotheses

- We can always look at a $(1-\alpha)100\%$ confidence interval and see what the result of a test would be if we carry out the test.
- For example, consider the $(1-\alpha)100\%$ confidence interval for μ

$$\left(\bar{y} - t_{\alpha/2} \frac{s}{\sqrt{n}}, \quad \bar{y} + t_{\alpha/2} \frac{s}{\sqrt{n}} \right)$$

- Suppose that μ_0 is **not included** in the above confidence interval because

$$\bar{y} - t_{\alpha/2} \frac{s}{\sqrt{n}} > \mu_0$$

Using Confidence Intervals to Test Hypotheses - cont.

- By rearranging this we see that this is equivalent to t_c being in the rejection region, i.e.,

$$\frac{\bar{y} - \mu_0}{s/\sqrt{n}} > t_{\alpha/2, (n-1)}$$

- Observe carefully that this is the same rejection region for the test of $H_0 : \mu \leq \mu_0$ vs. $H_a : \mu > \mu_0$ at level $\alpha/2$.
- That is, we will be rejecting H_0 at level $\alpha/2$ if

$$\frac{\bar{y} - \mu_0}{s/\sqrt{n}} > t_{\alpha/2, (n-1)}$$

- This is equivalent to saying that if μ_0 was not included in a $100(1-\alpha)\%$ interval, then H_0 will be rejected if the test is carried out at $\alpha/2$ level.

Using Confidence Intervals to Test Hypotheses - cont.

- For a **two-tailed test**, the confidence interval should be based on the same α as the test to make this inference.

In summary,

- To test $H_0 : \mu \leq \mu_0$, $H_a : \mu > \mu_0$, at level $\alpha/2$, use a $100(1-\alpha)\%$ confidence interval.
- To test $H_0 : \mu \geq \mu_0$, $H_a : \mu < \mu_0$, at level $\alpha/2$, use a $100(1-\alpha)\%$ confidence interval.
- $H_0 : \mu = \mu_0$, $H_a : \mu \neq \mu_0$, at level α , use a $100(1-\alpha)\%$ confidence interval.